

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT on

Machine Learning (23CS6PCMAL)

Submitted by

Arugunta Hamsika(1BM22CS054)

in partial fulfillment for the award of the degree of

**BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING**



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

Feb-2025 to June-2025

B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Arugunta Hamsika(1BM22CS054)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Sarala D V Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
---	--

Index

Sl. No.	Date	Experiment Title	Page No.
1	4-3-2025	Write a python program to import and export data using Pandas library functions	1-5
2	11-3-2025	Demonstrate various data pre-processing techniques for a given dataset	6-9
3	18-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	10-14
4	1-4-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	15-18
5	8-4-2025	Build Logistic Regression Model for a given dataset	19-23
6	15-4-2025	Build KNN Classification model for a given dataset.	24-30
7	15-4-2025	Build Support vector machine model for a given dataset	31-35
8	22-4-2025	Implement Random forest ensemble method on a given dataset.	36-39
9	22-4-2025	Implement Boosting ensemble method on a given dataset.	40-41
10	29-4-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	42-45
11	29-4-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	46-48

Github Link:

<https://github.com/hamsika04/6thSem-ML-La>

Program 1

Write a python program to import and export data using Pandas library functions

Code:

```
import pandas as pd
data = {
'Name': ['Alice', 'Bob', 'Charlie', 'David'],
'Age': [25, 30, 35, 40],
'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']
}
df = pd.DataFrame(data)
print("Sample data:")
print(df.head())
```

Sample data:

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago
3	David	40	Houston

```
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target
print("Sample data:")
print(df.head())
```

Sample data:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

target

0	0
1	0
2	0
3	0
4	0

```
from google.colab import files
uploaded = files.upload()
file_path = 'data.csv' # Ensure the file exists in the same directory
df = pd.read_csv(file_path)
```

```
print("Sample data:")
print(df.head())
print("\n")
```

Saving data.csv to data.csv

Sample data:

	ID	Name	Age	City
0	1	Alice	25	New York
1	2	Bob	30	Los Angeles
2	3	Charlie	35	Chicago
3	4	David	40	Houston
4	5	Eva	28	Phoenix

```
import yfinance as yf
import pandas as pd
import matplotlib.pyplot as plt
tickers = ["RELIANCE.NS", "TCS.NS", "INFY.NS"]
data = yf.download(tickers, start="2022-10-01", end="2023-10-01", group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
reliance_data = data['RELIANCE.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
```

First 5 rows of the dataset:

Ticker	RELIANCE.NS					
Price	Open	High	Low	Close	Volume	
Date						
2022-10-03	1092.199963	1103.822945	1079.184026	1082.152588	11852723	
2022-10-04	1095.077201	1104.302526	1091.583396	1102.110352	8948850	
2022-10-06	1109.326261	1118.916888	1104.371007	1106.174927	13352162	
2022-10-07	1102.772687	1116.131217	1102.772687	1110.856323	7714340	
2022-10-10	1098.365412	1104.119885	1090.601536	1098.730835	6329527	

Ticker	INFY.NS					
Price	Open	High	Low	Close	Volume	
Date						
2022-10-03	1337.743240	1337.743240	1313.110574	1320.453003	4943169	
2022-10-04	1345.038201	1356.928245	1339.638009	1354.228149	6631341	
2022-10-06	1369.007786	1383.029504	1368.155094	1378.624023	6180672	
2022-10-07	1370.286676	1381.181893	1364.412779	1374.881592	3994466	
2022-10-10	1351.338576	1387.956005	1351.338576	1385.729614	5274677	

Ticker	TCS.NS					
Price	Open	High	Low	Close	Volume	
Date						
2022-10-03	2799.044354	2823.062819	2779.418333	2789.651855	1763331	
2022-10-04	2831.707297	2895.304988	2825.212065	2888.903076	2145875	
2022-10-06	2907.454262	2919.603702	2890.117902	2898.996338	1790816	
2022-10-07	2894.744206	2901.847047	2858.015696	2864.370605	1939879	
2022-10-10	2813.062629	2922.407588	2808.389768	2914.510498	3064063	

Shape of the dataset:
(247, 15)

Column names:


```
MultiIndex([('RELIANCE.NS', 'Open'),
            ('RELIANCE.NS', 'High'),
            ('RELIANCE.NS', 'Low'),
            ('RELIANCE.NS', 'Close'),
            ('RELIANCE.NS', 'Volume'),
            ('INFY.NS', 'Open'),
            ('INFY.NS', 'High'),
            ('INFY.NS', 'Low'),
            ('INFY.NS', 'Close'),
            ('INFY.NS', 'Volume'),
            ('TCS.NS', 'Open'),
            ('TCS.NS', 'High'),
            ('TCS.NS', 'Low'),
            ('TCS.NS', 'Close'),
            ('TCS.NS', 'Volume')],
            names=['Ticker', 'Price'])
```

Summary statistics for Reliance Industries:

Price	Open	High	Low	Close	Volume
count	247.000000	247.000000	247.000000	247.000000	2.470000e+02
mean	1151.456593	1160.153640	1141.068147	1150.426972	1.316652e+07
std	66.114624	67.077801	65.976400	66.914372	6.754099e+06
min	1011.592400	1013.875900	995.607838	1005.312744	3.370033e+06
25%	1102.624229	1107.157058	1088.489391	1101.094238	8.717141e+06
50%	1151.342807	1158.969749	1142.665512	1151.160156	1.158959e+07
75%	1200.335361	1207.724151	1189.020599	1199.134460	1.530302e+07
max	1292.463484	1304.337734	1277.392357	1297.875366	5.708188e+07

```
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="Reliance Industries - Closing Price")
plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="Reliance Industries - Daily Returns", color='orange')
plt.tight_layout()
plt.show()
```

```
reliance_data.to_csv('reliance_stock_data.csv')
print("\nReliance stock data saved to 'reliance_stock_data.csv'.")
```



```
import yfinance as yf
import pandas as pd
import matplotlib.pyplot as plt
tickers = ["HDFCBANK.NS", "ICICIBANK.NS", "KOTAKBANK.NS"]
data = yf.download(tickers, start="2024-01-01", end="2024-12-30", group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
plt.figure(figsize=(14, 10))
for i, ticker in enumerate(tickers):
    bank_data = data[ticker]
    bank_data['Daily Return'] = bank_data['Close'].pct_change()
    plt.subplot(3, 2, 2*i+1)
    bank_data['Close'].plot(title=f'{ticker} - Closing Price', color='blue')
    plt.ylabel('Closing Price')
    plt.subplot(3, 2, 2*i+2)
    bank_data['Daily Return'].plot(title=f'{ticker} - Daily Returns', color='orange')
    plt.ylabel('Daily Return')
plt.tight_layout()
plt.show()
for ticker in tickers:
    bank_data = data[ticker]
    bank_data.to_csv(f'{ticker}_stock_data.csv')
    print(f'\n{ticker} stock data saved to '{ticker}_stock_data.csv'.")
```


first 5 rows of the dataset:

Ticker	Price	Open	High	Low	Close	Volume
2024-01-01	KOTAKBANK.NS	1006.909954	1916.899086	1891.027338	1067.859814	1425962
2024-01-02	KOTAKBANK.NS	1005.911188	1905.911188	1858.063525	1863.008179	5128796
2024-01-03	KOTAKBANK.NS	1061.959234	1867.952665	1845.627158	1863.857178	3781515
2024-01-04	KOTAKBANK.NS	1069.451068	1869.451068	1858.513105	1861.950602	2865768
2024-01-05	KOTAKBANK.NS	1063.457575	1867.852782	1839.383085	1845.577148	7799341


Ticker	Price	Open	High	Low	Close	Volume
2024-01-01	HDFCBANK.NS	1683.017598	1685.125187	1669.206199	1675.223999	7119843
2024-01-02	HDFCBANK.NS	1675.914685	1679.860799	1665.050651	1676.210571	14622046
2024-01-03	HDFCBANK.NS	1679.071480	1681.735059	1646.466666	1650.363525	14104881
2024-01-04	HDFCBANK.NS	1655.394910	1672.116520	1648.193203	1668.071777	13367028
2024-01-05	HDFCBANK.NS	1664.421596	1681.932477	1645.628188	1659.538288	15944735

Ticker	Price	Open	High	Low	Close	Volume
2024-01-01	ICICIBANK.NS	983.085778	996.273246	982.541485	990.869812	7683792
2024-01-02	ICICIBANK.NS	988.400253	989.134738	971.883221	973.866150	16263825
2024-01-03	ICICIBANK.NS	976.295294	979.567116	966.777197	975.650818	16826752
2024-01-04	ICICIBANK.NS	977.980767	980.707295	973.519176	978.724365	22789140
2024-01-05	ICICIBANK.NS	979.567084	989.779158	975.482920	985.218445	14875499

Shape of the dataset:
(244, 15)

Column names:

```
MultiIndex([('KOTAKBANK.NS', 'Open'),
              ('KOTAKBANK.NS', 'High'),
              ('KOTAKBANK.NS', 'Low'),
              ('KOTAKBANK.NS', 'Close'),
              ('KOTAKBANK.NS', 'Volume'),
              ('HDFCBANK.NS', 'Open'),
              ('HDFCBANK.NS', 'High'),
              ('HDFCBANK.NS', 'Low'),
              ('HDFCBANK.NS', 'Close'),
              ('HDFCBANK.NS', 'Volume'),
              ('ICICIBANK.NS', 'Open'),
              ('ICICIBANK.NS', 'High'),
              ('ICICIBANK.NS', 'Low'),
              ('ICICIBANK.NS', 'Close'),
              ('ICICIBANK.NS', 'Volume')],
            names=['Ticker', 'Price'])
```



HDFCBANK.NS stock data saved to 'HDFCBANK.NS_stock_data.csv'.

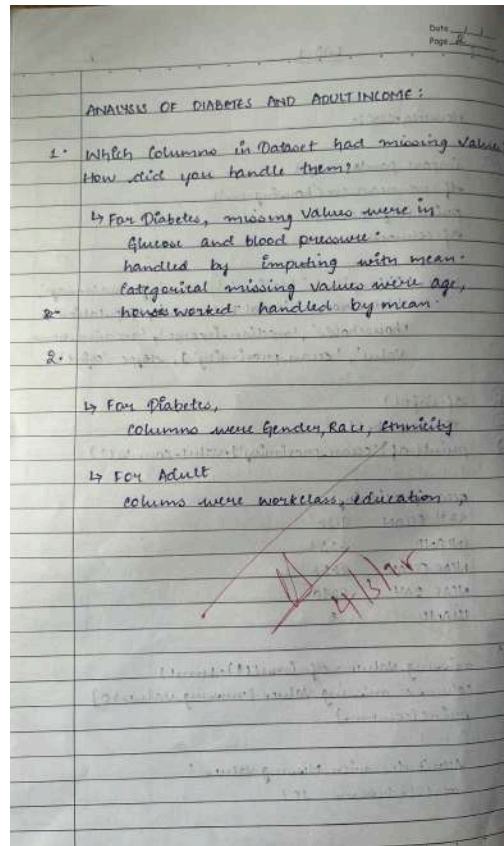
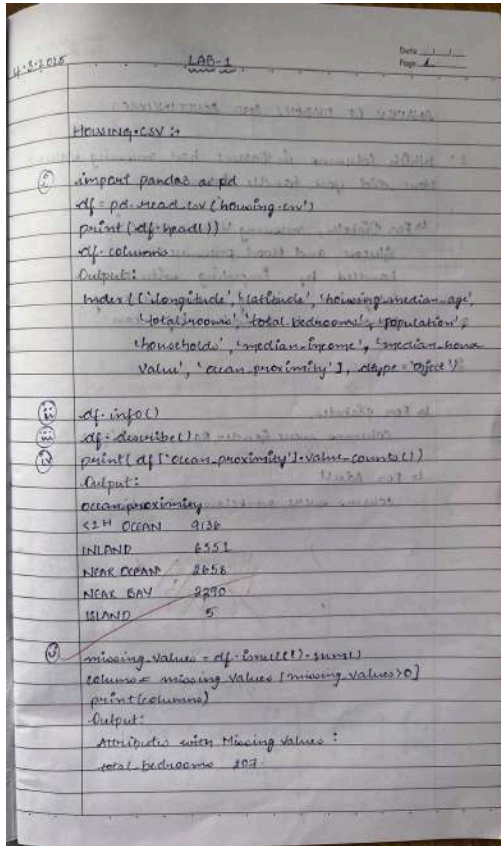
ICICIBANK.NS stock data saved to 'ICICIBANK.NS_stock_data.csv'.

KOTAKBANK.NS stock data saved to 'KOTAKBANK.NS_stock_data.csv'.

Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot:



Code:

```
import pandas as pd
file_path = 'housing.csv'
df = pd.read_csv(file_path)
print("\nDataset Information:")
print(df.info())
print("\nStatistical Information of Numerical Columns:")
print(df.describe())
print("\nUnique Labels Count for 'Ocean Proximity' column:")
print(df['ocean_proximity'].value_counts())
print("\nAttributes with Missing Values:")
missing_values = df.isnull().sum()
columns_with_missing_values = missing_values[missing_values > 0]
print(columns_with_missing_values)
```

```

Dataset Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28640 entries, 0 to 28639
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   longitude            28640 non-null  float64
1   latitude             28640 non-null  float64
2   housing_median_age   28640 non-null  float64
3   total_rooms          28640 non-null  float64
4   total_bedrooms       28433 non-null  float64
5   population            28640 non-null  float64
6   households            28640 non-null  float64
7   median_income        28640 non-null  float64
8   median_house_value   28640 non-null  float64
9   ocean_proximity      28640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
None

Statistical Information of Numerical Columns:
      longitude      latitude      housing_median_age      total_rooms \
count  28640.000000  28640.000000  28640.000000  28640.000000
mean   -119.569784    35.631861    28.639486   2635.763881
std      2.003532     2.135952    12.585558  2181.615252
min    -124.350000    32.540000     1.000000     2.000000
25%    -121.800000    33.930000    18.000000  1447.750000
50%    -118.400000    34.260000    29.000000  2127.000000
75%    -118.010000    37.710000    37.000000  3148.000000
max    -114.310000    41.950000    52.000000  39320.000000

      total_bedrooms      population      households      median_income \
count  28433.000000  28640.000000  28640.000000  28640.000000
mean     537.476744   3425.476744   499.539680     3.870671
std    421.385070   1132.462122   382.329753     1.899822
min       1.000000     3.000000     1.000000     0.499900
25%     295.000000   787.000000   280.000000     2.563400
50%     435.000000  1166.000000   409.000000     3.534800
75%     647.000000  1725.000000   605.000000     4.743250
max    6445.000000  35682.000000  6082.000000    15.000100

      median_house_value
count      28640.000000
mean    206855.816000
std    115395.615874
min     14999.000000
25%    119600.000000
50%    179700.000000
75%    264725.000000
max    500001.000000

Unique Labels Count for 'Ocean Proximity' column:
ocean_proximity
<IH OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND        5
Name: count, dtype: int64

Attributes with Missing Values:
total_bedrooms    207
dtype: int64

```

```

import pandas as pd
import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler, StandardScaler
file_path = "diabetes.csv"
df = pd.read_csv(file_path)
df_numeric = df.select_dtypes(include=['number']).copy() # Select only numeric columns
imputer = SimpleImputer(strategy="mean")
df_numeric.iloc[:, :] = imputer.fit_transform(df_numeric)
df[df_numeric.columns] = df_numeric
Q1 = df_numeric.quantile(0.25) # Only compute quartiles on numeric data
Q3 = df_numeric.quantile(0.75) # Only compute quartiles on numeric data
IQR = Q3 - Q1
df = df[~((df_numeric < (Q1 - 1.5 * IQR)) | (df_numeric > (Q3 + 1.5 * IQR))).any(axis=1)]
min_max_scaler = MinMaxScaler()
df_minmax = pd.DataFrame(min_max_scaler.fit_transform(df_numeric),
columns=df_numeric.columns) # Only transform the numeric columns
standard_scaler = StandardScaler()
df_standard = pd.DataFrame(standard_scaler.fit_transform(df_numeric), columns=df_numeric.columns)
# Only transform the numeric columns
print("\nProcessed Diabetes Dataset (Min-Max Scaled):")
print(df_minmax.head())

```

```
print("\nProcessed Diabetes Dataset (Standard Scaled):")
print(df_standard.head())
```

```
Processed Diabetes Dataset (Min-Max Scaled):
```

	ID	No_Patien	AGE	Urea	Cr	HbA1c	Chol	\
0	0.627034	0.000237	0.508475	0.109375	0.050378	0.264901	0.407767	
1	0.918648	0.000452	0.101695	0.104167	0.070529	0.264901	0.359223	
2	0.524406	0.000634	0.508475	0.109375	0.050378	0.264901	0.407767	
3	0.849812	0.001160	0.508475	0.109375	0.050378	0.264901	0.407767	
4	0.629537	0.000452	0.220339	0.171875	0.050378	0.264901	0.475728	

	TG	HDL	LDL	VLDL	BMI
0	0.044444	0.226804	0.114583	0.011461	0.173913
1	0.081481	0.092784	0.187500	0.014327	0.139130
2	0.044444	0.226804	0.114583	0.011461	0.173913
3	0.044444	0.226804	0.114583	0.011461	0.173913
4	0.051852	0.061856	0.177083	0.008596	0.069565


```
Processed Diabetes Dataset (Standard Scaled):
```

	ID	No_Patien	AGE	Urea	Cr	HbA1c	Chol	\
0	0.672140	-0.074747	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	
1	1.641852	-0.069940	-3.130017	-0.212954	-0.115804	-1.334983	-0.893730	
2	0.330868	-0.065869	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	
3	1.412950	-0.054126	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	
4	0.680463	-0.069939	-2.334096	0.673299	-0.382672	-1.334983	0.028576	

	TG	HDL	LDL	VLDL	BMI
0	-1.035084	1.810756	-1.085457	-0.369958	-1.124622
1	-0.678063	-0.158692	-0.457398	-0.342649	-1.326239
2	-1.035084	1.810756	-1.085457	-0.369958	-1.124622
3	-1.035084	1.810756	-1.085457	-0.369958	-1.124622
4	-0.963680	-0.613180	-0.547121	-0.397267	-1.729472

```
import pandas as pd
import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder
file_path = "adult.csv" # Update the path if needed
df = pd.read_csv(file_path)
df.replace("?", np.nan, inplace=True)
num_imputer = SimpleImputer(strategy="mean")
df[df.select_dtypes(include=['number']).columns] =
num_imputer.fit_transform(df.select_dtypes(include=['number']))
cat_imputer = SimpleImputer(strategy="most_frequent")
df[df.select_dtypes(include=['object']).columns] =
cat_imputer.fit_transform(df.select_dtypes(include=['object']))
label_encoders = {}
for col in df.select_dtypes(include=['object']).columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1
```

```

df = df[~((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).any(axis=1)]
min_max_scaler = MinMaxScaler()
df_minmax = pd.DataFrame(min_max_scaler.fit_transform(df), columns=df.columns)
standard_scaler = StandardScaler()
df_standard = pd.DataFrame(standard_scaler.fit_transform(df), columns=df.columns)
print("\nProcessed Adult Income Dataset (Min-Max Scaled):")
print(df_minmax.head())
print("\nProcessed Adult Income Dataset (Standard Scaled):")
print(df_standard.head())

```

Processed Adult Income Dataset (Min-Max Scaled):

	age	workclass	fnlwgt	education	educational-num	marital-status	\
0	0.344262	0.0	0.188277	0.555556	0.363636	0.333333	
1	0.114754	0.0	0.881156	1.000000	0.454545	0.666667	
2	0.147541	0.0	0.169156	0.555556	0.363636	0.666667	
3	0.672131	0.0	0.708251	0.555556	0.363636	0.333333	
4	0.131148	0.0	0.475807	0.333333	0.727273	0.333333	

	occupation	relationship	race	gender	capital-gain	capital-loss	\
0	0.307692	0.0	0.0	1.0	0.0	0.0	
1	0.538462	0.8	0.0	0.0	0.0	0.0	
2	0.000000	0.2	0.0	0.0	0.0	0.0	
3	0.692308	0.0	0.0	1.0	0.0	0.0	
4	0.692308	0.0	0.0	1.0	0.0	0.0	

	hours-per-week	native-country	income
0	0.894737	0.0	0.0
1	0.368421	0.0	0.0
2	0.315789	0.0	0.0
3	0.105263	0.0	0.0
4	0.368421	0.0	0.0

Processed Adult Income Dataset (Standard Scaled):

	age	workclass	fnlwgt	education	educational-num	marital-status	\
0	0.220179	0.0	-1.022983	-0.151256	-0.654083	-0.398228	
1	-0.955630	0.0	2.234629	1.457372	-0.073261	0.828047	
2	-0.787657	0.0	-1.112882	-0.151256	-0.654083	0.828047	
3	1.899906	0.0	1.421707	-0.151256	-0.654083	-0.398228	
4	-0.871644	0.0	0.328856	-0.955571	1.669207	-0.398228	

	occupation	relationship	race	gender	capital-gain	capital-loss	\
0	-0.420679	-1.044582	0.0	0.770972	0.0	0.0	
1	0.305840	1.629927	0.0	-1.297064	0.0	0.0	
2	-1.389371	-0.375955	0.0	-1.297064	0.0	0.0	
3	0.790186	-1.044582	0.0	0.770972	0.0	0.0	
4	0.790186	-1.044582	0.0	0.770972	0.0	0.0	

	hours-per-week	native-country	income
0	2.312838	0.0	0.0
1	-0.329781	0.0	0.0
2	-0.594043	0.0	0.0
3	-1.651090	0.0	0.0
4	-0.329781	0.0	0.0

Program 3

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot:

LAB-11 (Decision Tree)

Question 1:

Instance	a_1	a_2	Classification
1	hot	high	No
2	hot	high	No
6	cool	high	No
7	hot	high	No
8	hot	normal	Yes

Solution:

Entropy(S) = $-\frac{4}{8} \log_2 \left(\frac{4}{8}\right) - \frac{4}{8} \log_2 \left(\frac{4}{8}\right)$
 $= -0.5 \log_2 \left(\frac{1}{2}\right) - 0.5 \log_2 \left(\frac{1}{2}\right)$
 $= 0.7219$

For a_2 :

Hot: $\{1, 2, 7\}$ = $-\frac{1}{4} \log_2 \left(\frac{1}{4}\right) - \frac{3}{4} \log_2 \left(\frac{3}{4}\right)$
 $= 0.8113$

Cool: $\{6\}$ = 0

Gain(a_2) = $0.7219 - \frac{4}{8} \times 0.8113 - \frac{4}{8} \times 0$
 $= 0.3286$

For a_1 :

High: $\{1, 2, 6, 7\}$ = 0

Normal: $\{8\}$ = 0

Gain(a_1) = $0.7219 - 0 - 0 = 0.7219$

a_1 is the root. Since Gain(a_1) is high

```

graph TD
    A((a1)) -- high --> B[NO]
    A -- normal --> C[YES]
    B --- D["1, 2, 6, 7"]
    C --- E["8"]
  
```

1. IRIS.csv

The accuracy score for Iris Dataset was 1.0 meaning it correctly classified all test samples.

2. Confusion Matrix:

10	0	0
0	9	0
0	0	11

→ all 10 samples are classified as setosa
 → all 9 samples are classified as Versicolour
 → all 11 samples are classified as Virginica

→ Decision Tree is perfectly classified and there were no misclassifications.

CODE:

```

import pandas as pd
import numpy as np
iris_data = pd.read_csv("iris.csv")
X_iris = iris_data.iloc[:, 2:4]
y_iris = iris_data.iloc[:, 4:5]

X_train, X_test, y_train, y_test = train_test_split(X_iris, y_iris,
                                                    test_size=0.2, RandomState=42)

iris_classifier = DecisionTreeClassifier()
iris_classifier.fit(X_train, y_train)
y_pred_iris = iris_classifier.predict(X_test)

print("IRIS DATASET accuracy score (y_test, y_pred)")
print("IRIS DATASET CONFUSION MATRIX")
  
```

Output: Iris Dataset Accuracy: 1.0
 Iris Dataset Confusion Matrix:

10	0	0
0	9	0
0	0	11

Code:

Iris.csv

```

import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv('iris.csv')
le = LabelEncoder()
df['species'] = le.fit_transform(df['species'])
X = df.drop('species', axis=1)
y = df['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
clf = DecisionTreeClassifier(criterion='entropy', random_state=42)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
  
```

```

accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
print(classification_report(y_test, y_pred, target_names=le.classes_))
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=le.classes_,
yticklabels=le.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt
plt.figure(figsize=(12,8))
plot_tree(clf, filled=True, feature_names=X.columns)
plt.show()

```


```

Accuracy: 1.00
          precision    recall  f1-score   support

 Iris-setosa         1.00      1.00      1.00        10
 Iris-versicolor     1.00      1.00      1.00         9
 Iris-virginica       1.00      1.00      1.00        11

 accuracy               1.00      1.00      1.00        30
 macro avg              1.00      1.00      1.00        30
 weighted avg           1.00      1.00      1.00        30

```




Drug.csv

```
import pandas as pd
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
df = pd.read_csv('drug.csv')
label_encoders = {}
for column in df.columns:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le
X = df.drop('Drug', axis=1)
y = df['Drug']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
clf = DecisionTreeClassifier(criterion='entropy')
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

# Evaluate the classifier
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
print(classification_report(y_test, y_pred))
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6,4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=le.classes_,
yticklabels=le.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt
plt.figure(figsize=(12,8))
plot_tree(clf, filled=True, feature_names=X.columns)
plt.show()
```


Accuracy: 1.00

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6
1	1.00	1.00	1.00	3
2	1.00	1.00	1.00	5
3	1.00	1.00	1.00	11
4	1.00	1.00	1.00	15
accuracy			1.00	40
macro avg	1.00	1.00	1.00	40
weighted avg	1.00	1.00	1.00	40



```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.metrics import mean_absolute_error, mean_squared_error
import matplotlib.pyplot as plt
import math

data = pd.read_csv('petrol_consumption.csv')
X = data.drop('Petrol_Consumption', axis=1)
y = data['Petrol_Consumption']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
regressor = DecisionTreeRegressor(random_state=42, max_depth=3) # Limiting depth for readability
regressor.fit(X_train, y_train)
y_pred = regressor.predict(X_test)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = math.sqrt(mse)
```

```
print("Regression Tree Evaluation Metrics:")
print(f'Mean Absolute Error (MAE): {mae:.2f}')
print(f'Mean Squared Error (MSE): {mse:.2f}')
print(f'Root Mean Squared Error (RMSE): {rmse:.2f}')
plt.figure(figsize=(20,6))
plot_tree(
    regressor,
    feature_names=X.columns,
    filled=True,
    rounded=True,
    fontsize=10,
)
plt.title("Regression Tree for Petrol Consumption Prediction", fontsize=14)
plt.show()
```

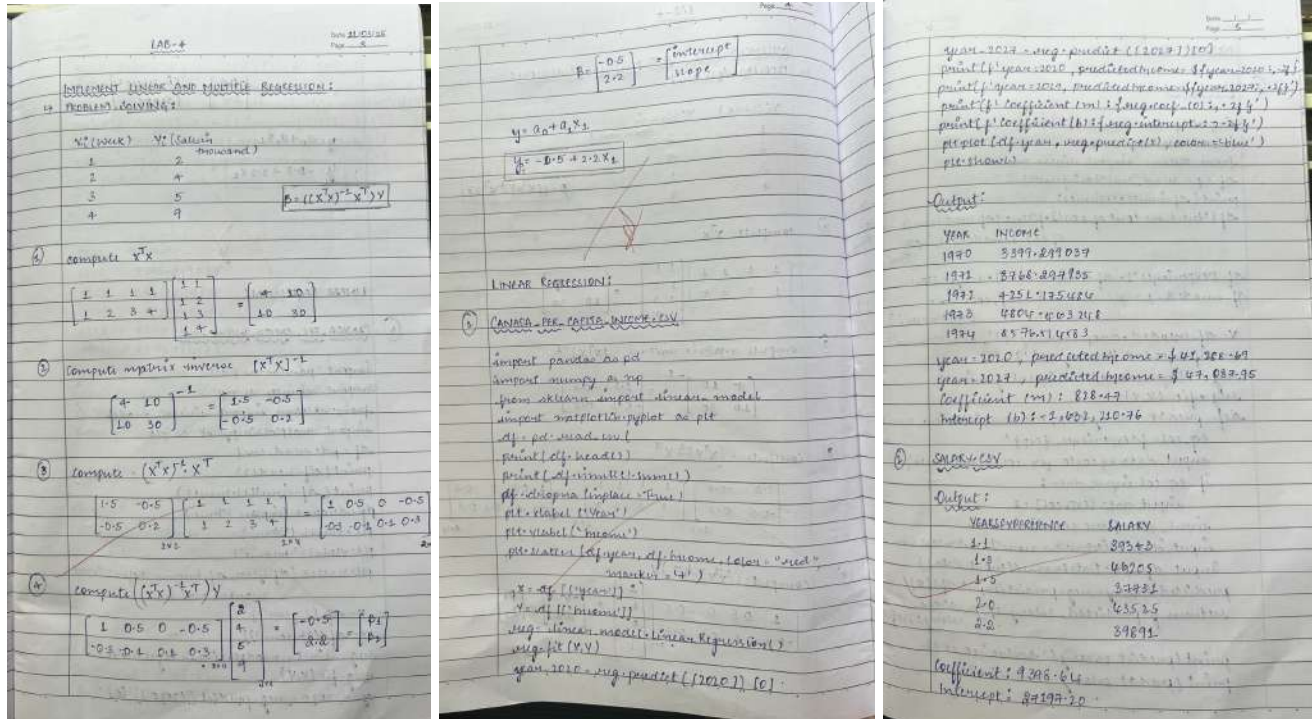
```
Regression Tree Evaluation Metrics:
Mean Absolute Error (MAE): 80.63
Mean Squared Error (MSE): 14718.40
Root Mean Squared Error (RMSE): 121.32
```

Program 4

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset.

Linear Regression:

Screenshot:




Code:

```
import pandas as pd
import numpy as np
from sklearn import linear_model
import matplotlib.pyplot as plt
df = pd.read_csv('canada_per_capita_income.csv')
reg = linear_model.LinearRegression()
reg.fit(df[['year']], df['per capita income (US$)'])
print(f'Coefficient: {reg.coef_}')
print(f'Intercept: {reg.intercept_}')
predicted_income = reg.predict([[2020]])
print(f'Predicted per capita income for 2020: ${predicted_income[0]:.2f}')
plt.scatter(df['year'], df['per capita income (US$)'], color='red')
plt.plot(df['year'], reg.predict(df[['year']]), color='blue')
plt.xlabel('Year')
plt.ylabel('Per Capita Income (US$)')
plt.show()
print(results.head())
```

Coefficient: [828.46507522]

Intercept: -1632210.7578554575

Predicted per capita income for 2020: \$41,288.69



Screenshot:

```

Date: _____
Page: ____ of ____

MULTIPLE REGRESSION:

① Hiring.csv

import pandas as pd
import numpy as np
from sklearn import linear_model
df = pd.read_csv('Hiring.csv')
print(df.shape) # (rows, columns)
df['test score (out of 10)'] = df['test score (out of 10)'].fillna(df['test score (out of 10)'].max())

df.experience = df['experience'].astype(int)
df['enacted'] = pd.get_dummies(df, column = ['experience'])

X = df.exacted.drop('Salary', 1).axis('columns')
y = df.exacted['Salary']

reg = linear_model.LinearRegression()
reg.fit(X, y)

def predict_salary (exp, test score, interview score):
    exp_col = 'experience'
    input_data = {}
    for col in X.columns:
        if col != exp_col:
            input_data[col] = test score
        else:
            input_data[col] = exp

    input_data['interview score'] = interview score
    input_data['test score'] = test score
    input_df = pd.DataFrame([input_data])
    predicted_salary = reg.predict(input_df)[0]
    return f'predicted salary is {predicted_salary}'

print(predict_salary('thorvald', 10, 10))
print(predict_salary('Thorvald', 10, 10))

```


Output:

experience	2
test score	1
Interview score	0
Salary (\$)	0

predicted salary: \$ 67,000.99
predicted salary: \$ 63,869.80

① COMPANIES = CSV

predicted profit: base薪=0.99



Code:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
df_hiring = pd.read_csv('hiring.csv')
print("Original Data:")
print(df_hiring.head())
df_hiring['experience'] = df_hiring['experience'].replace({
    'five': 5,
    'four': 4,
    'three': 3
})
df_hiring['experience'] = pd.to_numeric(df_hiring['experience'], errors='coerce')
df_hiring['test_score(out of 10)'] = pd.to_numeric(df_hiring['test_score(out of 10)'], errors='coerce')
df_hiring['interview_score(out of 10)'] = pd.to_numeric(df_hiring['interview_score(out of 10)'],
errors='coerce')
df_hiring['salary($)'] = pd.to_numeric(df_hiring['salary($)'], errors='coerce')
df_hiring['experience'] = df_hiring['experience'].fillna(df_hiring['experience'].median())
df_hiring['test_score(out of 10)'] = df_hiring['test_score(out of 10)'].fillna(df_hiring['test_score(out of
10)'].median())
df_hiring['interview_score(out of 10)'] = df_hiring['interview_score(out of
10)'].fillna(df_hiring['interview_score(out of 10)'].median())
df_hiring['salary($)'] = df_hiring['salary($)'].fillna(df_hiring['salary($)'].median())
X_hiring = df_hiring[['experience', 'test_score(out of 10)', 'interview_score(out of 10)']]
y_hiring = df_hiring['salary($)']
```

```

model_hiring = LinearRegression()
model_hiring.fit(X_hiring, y_hiring)
print(f"\nCoefficients: {model_hiring.coef_}")
print(f"Intercept: {model_hiring.intercept_}")
predictions_hiring = model_hiring.predict([[2, 9, 6], [12, 10, 10]])
print(f"\nPredicted salary for candidate 1 (2 years experience, 9 test score, 6 interview score):
${predictions_hiring[0]:,.2f}")
print(f"Predicted salary for candidate 2 (12 years experience, 10 test score, 10 interview score):
${predictions_hiring[1]:,.2f}")

```

Original Data:

	experience	test_score(out of 10)	interview_score(out of 10)	salary(\$)
0	NaN	8.0	9	50000
1	NaN	8.0	6	45000
2	five	6.0	7	60000
3	two	10.0	10	65000
4	seven	9.0	6	70000

```

Coefficients: [-793.62416107  -5.03355705  139.26174497]
Intercept: 65117.44966442955

```

Predicted salary for candidate 1 (2 years experience, 9 test score, 6 interview score): \$64,320.47

Predicted salary for candidate 2 (12 years experience, 10 test score, 10 interview score): \$56,936.24

Program 5

Build Logistic Regression Model for a given dataset.

Binary Logical Regression: Screenshot:

LAB-3

(i) Binary Classification Problem:
Given,
 $a_0 = -5$
 $a_1 = 0.8$
 $x = 7$
 $z = a_0 + a_1 x$
 $= -5 + (0.8 \times 7)$
 $= 0.6$
 $y = \frac{1}{1 + e^{-z}}$
 $= \frac{1}{1 + e^{-0.6}}$
 $= 0.64$
Given threshold = 0.5
 $0.64 > 0.5$
∴ Student will study for 7 hours will pass.

(ii) SOFTMAX FUNCTION:
 $z = [a_0, a_1, a_2]$ for 3 classes
 $z_1 = \frac{e^{z_1}}{e^{z_1} + e^{z_2} + e^{z_3}}$
 $z_1 = \frac{e^{-3.5}}{e^{-3.5} + e^{-2.5} + e^{-1.5}} = 0.665$
 $z_2 = \frac{e^{-2.5}}{e^{-3.5} + e^{-2.5} + e^{-1.5}} = 0.244$

LOGISTIC REGRESSION MULTICLASS:

```
CODE-  
import pandas as pd  
from sklearn.datasets import load_iris  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import accuracy_score  
from sklearn.metrics import confusion_matrix  
import matplotlib.pyplot as plt  
iris = load_iris(return_X_y=True)  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size=0.2, random_state=0)  
model = LogisticRegression(multi_class='multinomial')  
model.fit(X_train, y_train)  
y_pred = model.predict(X_test)  
accuracy = accuracy_score(y_test, y_pred)  
confusion_matrix = confusion_matrix(y_test, y_pred)  
cm = display.plot()
```

Output:

Accuracy of the Multinomial Logistic Regression Model on the test set is 1.00

Actual \ Predicted	Setosa	Versicolour	Virginica
Setosa	10	0	0
Versicolour	0	9	0
Virginica	0	0	11

Predicted label

LOGISTIC REGRESSION BINARY:

```
CODE-  
import pandas as pd  
from sklearn.datasets import load_iris  
X = pd.read_csv('iris_data.csv')  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
                                                    test_size=0.2, random_state=0)  
model = LogisticRegression()  
model.fit(X_train, y_train)  
y_pred = model.predict(X_test)  
accuracy = accuracy_score(y_test, y_pred)  
confusion_matrix = confusion_matrix(y_test, y_pred)
```

```
X_test  
y_test  
y_pred = model.predict(X_test)  
y_pred  
model.score(X_test, y_test)  
model.predict_proba(X_test)  
y_pred = model.predict(X_test)  
y_pred  
model.coef_  
model.intercept_  
import math  
def sigmoid(x):  
    return 1 / (1 + math.exp(-x))  
def prediction_function(age):  
    x = 0.13 * age - 4.988  
    y = sigmoid(x)  
    return y  
age = 35  
prediction_function(age)  
0.57 is less than 0.5 which means person with  
35 will not buy insurance  
SAMPLE PROGRAM QUESTIONS:  
1. HR comma, as per cv  
2. Salary has direct impact on employee retention  
because as per the data employees who  
have less salary have a lower retention  
rate & vice versa
```

(ii) The accuracy of logistic Regression model is 76%. The accuracy is decent. (Neither good nor bad)

2. Zoo Dataset:

(i) I have dropped the columns 'animal-name' and 'class-type' as they have categorical value.

(ii) there weren't any missing values in dataset.

(iii) the Confusion matrix tell us that there are higher number of positive values.

(iv) Reptile class has been misclassified as Fish class.

Code:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
hr_data = pd.read_csv("HR_comma_sep.csv")
print(hr_data.head())
print("\nMissing values:\n", hr_data.isnull().sum())
print("\nData Types:\n", hr_data.dtypes)
print("\nUnique values in categorical columns:\n",
hr_data.select_dtypes(include=['object']).nunique())
plt.figure(figsize=(8,5))
sns.countplot(x="salary", hue="left", data=hr_data)
plt.title("Impact of Salary on Employee Retention")
plt.xlabel("Salary Level")
plt.ylabel("Number of Employees")
plt.legend(["Stayed", "Left"])
plt.show()
```


	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	

	time_spend_company	Work_accident	left	promotion_last_5years	Department	\
0	3	0	1	0	sales	
1	6	0	1	0	sales	
2	4	0	1	0	sales	
3	5	0	1	0	sales	
4	3	0	1	0	sales	

	salary
0	low
1	medium
2	medium
3	low
4	low

Missing values:

satisfaction_level	0
last_evaluation	0
number_project	0
average_monthly_hours	0
time_spend_company	0
Work_accident	0
left	0
promotion_last_5years	0
Department	0
salary	0

dtype: int64

Data Types:


satisfaction_level	float64
last_evaluation	float64
number_project	int64
average_monthly_hours	int64
time_spend_company	int64
Work_accident	int64
left	int64
promotion_last_5years	int64
Department	object
salary	object

dtype: object

Unique values in categorical columns:

Department	10
salary	3

dtype: int64



Multi Linear classification

Code:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
zoo_data = pd.read_csv("zoo-data.csv")
class_types = pd.read_csv("zoo-class-type.csv")
print("Zoo Data:\n", zoo_data.head())
print("\nClass Type Data:\n", class_types.head())
if "class_type" not in zoo_data.columns:
    print("\nError: 'class_type' column is missing in zoo-data.csv. Available columns: ",
zoo_data.columns)
else:
    X = zoo_data.drop(columns=["animal_name", "class_type"], errors="ignore")
    y = zoo_data["class_type"]
    print("\nUnique class types:", y.unique())
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
    scaler = StandardScaler()
    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)
    model = LogisticRegression(multi_class="multinomial", solver="lbfgs", max_iter=200)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"\nModel Accuracy: {accuracy:.4f}")
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(8,6))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=y.unique(), yticklabels=y.unique())
    plt.xlabel("Predicted Class")
    plt.ylabel("Actual Class")
    plt.title("Confusion Matrix - Zoo Dataset")
    plt.show()
    print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```

Zoo Data:
  animal_name  hair  feathers  eggs  milk  airborne  aquatic  predator  \
0   aardvark    1        0      0     1        0        0         1
1   antelope    1        0      0     1        0        0         0
2     bass     0        0      1     0        0        1         1
3     bear     1        0      0     1        0        0         1
4     boar     1        0      0     1        0        0         1

  toothed  backbone  breathes  venomous  fins  legs  tail  domestic  catsize  \
0         1         1         1         0     0    4     0         0         1
1         1         1         1         0     0    4     1         0         1
2         1         1         0         0     1    0     1         0         0
3         1         1         1         0     0    4     0         0         1
4         1         1         1         0     0    4     1         0         1

  class_type
0           1
1           1
2           4
3           1
4           1

```

```

Class Type Data:
  Class_Number  Number_Of_Animal_Species_In_Class  Class_Type  \
0             1                               41      Mammal
1             2                               20       Bird
2             3                               5       Reptile
3             4                               13       Fish
4             5                               4  Amphibian

```


```

Animal Names
0  aardvark, antelope, bear, boar, buffalo, calf,...
1  chicken, crow, dove, duck, flamingo, gull, haw...
2  pitviper, seasnake, slowworm, tortoise, tuatara
3  bass, carp, catfish, chub, dogfish, haddock, h...
4  frog, frog, newt, toad

```

Unique class types: [1 4 2 7 6 5 3]

Model Accuracy: 0.9524



```

Classification Report:
      precision    recall  f1-score   support

     1         1.00      1.00      1.00        12
     2         1.00      1.00      1.00         2
     3         0.00      0.00      0.00         1
     4         1.00      1.00      1.00         2
     6         1.00      1.00      1.00         3
     7         0.50      1.00      0.67         1

 accuracy          0.75      0.83      0.95        21
 macro avg          0.75      0.83      0.78        21
 weighted avg          0.93      0.95      0.94        21

```

Program 6

Build KNN Classification model for a given dataset.

Screenshot:

LAB-6 (KNN)

Date: / / Page: /

Succession

Person	Age	Salary	Target	Distance	Rank
A	18	50	N	52.8	5
B	23	55	N	46.5	4
C	24	70	N	31.95	2
D	42	60	Y	40.44	3
E	43	70	Y	31.04	1
F	38	40	Y	60.07	6
X	35	100	Y		

K=3

→ for (35, 100) → 1st 43 70 Y
2nd 24 70 N
3rd 42 60 Y

(30, target for (35, 100) is Y.

CODE:

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

from sklearn.model_selection import train_test_split
import pandas as pd

iris_data = pd.read_csv('iris.csv')
X_train, X_test, y_train, y_test = train_test_split(
    iris_data[['sepal_length', 'sepal_width', 'petal_length', 'petal_width']],
    iris_data['species'], test_size=0.2, random_state=42)

knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)

```

LAB-6 (KNN)

Date: / / Page: /

Succession

```

knn_classifier = KNeighborsClassifier(n_neighbors=3)
y_pred_test = knn_classifier.predict(X_test)

print("Accuracy Score: ", accuracy_score(y_test, y_pred_test))

print("Confusion Matrix: ", confusion_matrix(y_test, y_pred_test))

print("Classification Report: ")

```

Output:

Accuracy Score: 1.0

Confusion Matrix:

10	0	0
0	9	0
0	0	11

Classification Report:

	precision	recall	f1-score	support
setosa	1	1	1	10
versicolour	1	1	1	9
virginica	1	1	1	11
accuracy			1	30

IRIS DATASET:

to choose best K value, you can test with different K values. Check accuracy rate and error rate. small K might give overfit data but large K can underfit data. so in this case K=3 gives perfect accuracy.

LAB-6 (KNN)

Date: / / Page: /

Succession

DIABETES DATASET:

- Feature Scaling is essential for pre-processing especially for algo like KNN which rely on distance calculations.
- Feature are equal importance, improved convergence and avoid distance bias.
- two methods to perform this:
 - 1) Min-Max Scaling
 - 2) Standardization $\rightarrow Z = \frac{X - \mu}{\sigma}$

Code:

Iris.csv

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (accuracy_score, confusion_matrix, classification_report,
ConfusionMatrixDisplay)
from sklearn.preprocessing import LabelEncoder


data = pd.read_csv('iris.csv')
X = data.drop('species', axis=1)
y = data['species']
le = LabelEncoder()
y = le.fit_transform(y)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
accuracy = []
for k in range(1, 21):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy.append(accuracy_score(y_test, y_pred))
plt.figure(figsize=(10, 6))

```

```

plt.plot(range(1, 21), accuracy, marker='o')
plt.title('Accuracy vs K Value')
plt.xlabel('K Value')
plt.ylabel('Accuracy')
plt.xticks(range(1, 21))
plt.grid()
plt.show()
optimal_k = np.argmax(accuracy) + 1 # +1 because range starts at 1
print(f"\nOptimal K value: {optimal_k}")
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"\nAccuracy: {accuracy:.4f}")
cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(cm)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=le.classes_)
disp.plot(cmap='Blues')
plt.title('Confusion Matrix')
plt.show()
print("\nClassification Report:")
print(classification_report(y_test, y_pred, target_names=le.classes_))

```




Optimal K value: 1

Accuracy: 1.0000

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



Classification Report:

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	10
Iris-versicolor	1.00	1.00	1.00	9
Iris-virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

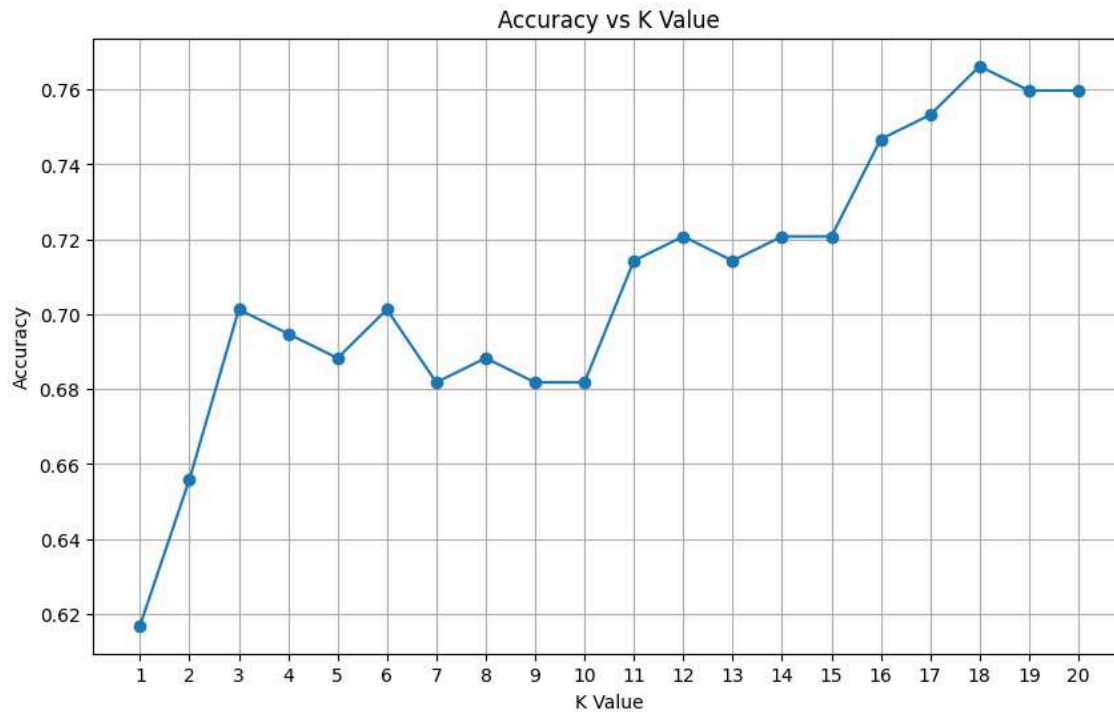
Diabetes.csv

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix, ConfusionMatrixDisplay
data = pd.read_csv('diabetes.csv')
X = data.drop('Outcome', axis=1)
y = data['Outcome']
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
k_values = range(1, 21)
accuracy_scores = []
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy_scores.append(accuracy_score(y_test, y_pred))
```

```

plt.figure(figsize=(10, 6))
plt.plot(k_values, accuracy_scores, marker='o')
plt.title('Accuracy vs K Value')
plt.xlabel('K Value')
plt.ylabel('Accuracy')
plt.xticks(k_values)
plt.grid()
plt.show()
optimal_k = k_values[np.argmax(accuracy_scores)]
print(f'Optimal K value: {optimal_k}')
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'\nAccuracy: {accuracy:.4f}')
cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(cm)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['No Diabetes', 'Diabetes'])
disp.plot(cmap='Blues')
plt.title('Confusion Matrix')
plt.show()
print("\nModel Evaluation:")
print(f'- True Positives: {cm[1,1]}')
print(f'- True Negatives: {cm[0,0]}')
print(f'- False Positives: {cm[0,1]}')
print(f'- False Negatives: {cm[1,0]}')

```




Optimal K value: 18

Accuracy: 0.7662

Confusion Matrix:

```
[[89 10]
 [26 29]]
```



Model Evaluation:

- True Positives: 29
- True Negatives: 89
- False Positives: 10
- False Negatives: 26


Heart.csv

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import cross_val_score
try:
    df = pd.read_csv("heart.csv")
except FileNotFoundError:
    print("Error: 'heart.csv' not found. Please make sure the file is in the correct directory.")
    exit()
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
k_values = list(range(1, 31, 2)) # Try odd k values from 1 to 30
cv_scores = []
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train_scaled, y_train, cv=5, scoring='accuracy')
    cv_scores.append(scores.mean())
plt.figure(figsize=(12, 6))
plt.plot(k_values, cv_scores, marker='o')
plt.title('Cross-Validation Accuracy vs. K Value (Scaled Data)')
plt.xlabel('Number of Neighbors (K)')
plt.ylabel('Mean Cross-Validation Accuracy')
plt.xticks(k_values)
plt.grid(True)
plt.show()
best_k_index = cv_scores.index(max(cv_scores))
best_k = k_values[best_k_index]
print(f"\nThe optimal K value found through cross-validation is: {best_k}")
knn_classifier = KNeighborsClassifier(n_neighbors=best_k)
knn_classifier.fit(X_train_scaled, y_train)
y_pred = knn_classifier.predict(X_test_scaled)
accuracy = accuracy_score(y_test, y_pred)
confusion = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred)
print(f"\nK-Nearest Neighbors Classifier with K = {best_k} (Scaled Data)")
print("Accuracy Score on Test Data:", accuracy)
plt.figure(figsize=(8, 6))
sns.heatmap(confusion, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Heart Disease', 'Heart Disease'], yticklabels=['No Heart Disease', 'Heart
```

```

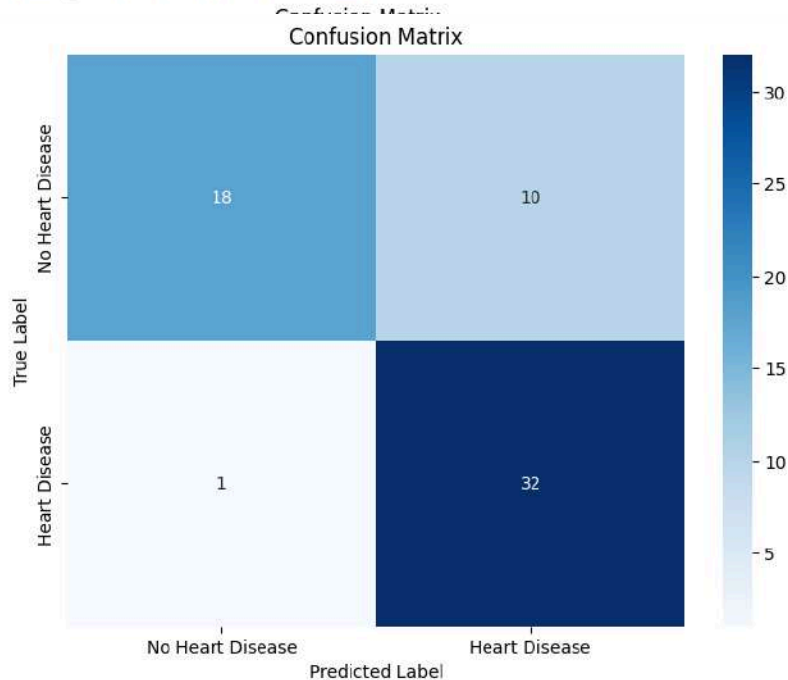
Disease'])
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
print("\nClassification Report on Test Data:\n", report)

```



The optimal K value found through cross-validation is: 21

K-Nearest Neighbors Classifier with K = 21 (Scaled Data)
Accuracy Score on Test Data: 0.819672131147541

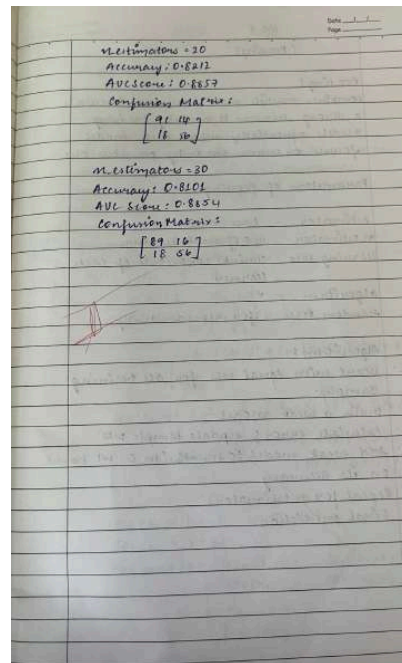
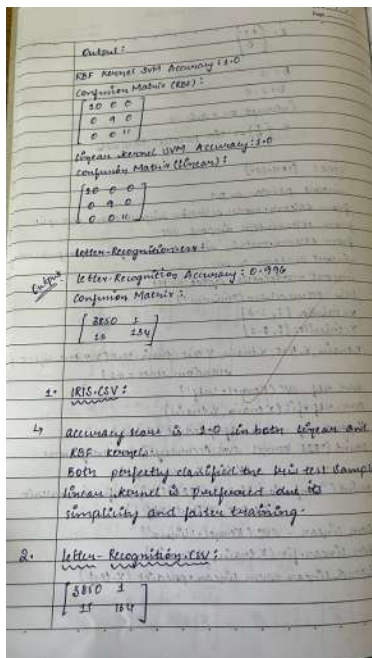
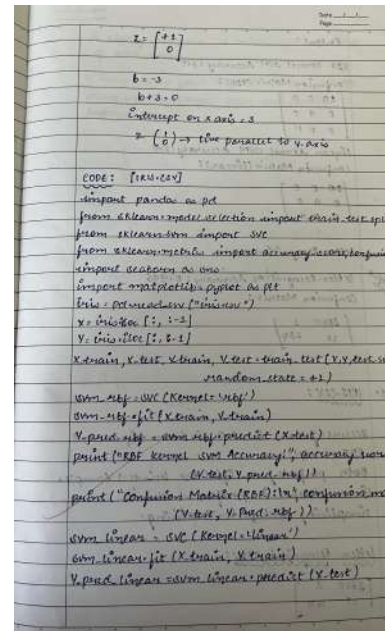
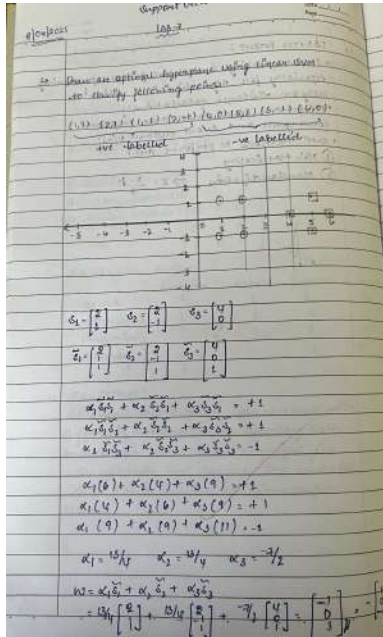


Classification Report on Test Data:

	precision	recall	f1-score	support
0	0.95	0.64	0.77	28
1	0.76	0.97	0.85	33
accuracy			0.82	61
macro avg	0.85	0.81	0.81	61
weighted avg	0.85	0.82	0.81	61

Program 7

Build Support vector machine model for a given dataset
Screenshot:



Code:


Iris.csv

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```


```

iris_df = pd.read_csv('iris.csv')
X = iris_df.iloc[:, :-1]
y = iris_df.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
def evaluate_svm(kernel_type):
    svm_clf = SVC(kernel=kernel_type, random_state=42)
    svm_clf.fit(X_train, y_train)
    y_pred = svm_clf.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(6, 4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=svm_clf.classes_,
                yticklabels=svm_clf.classes_)
    plt.title(f'Confusion Matrix ({kernel_type} kernel)')
    plt.ylabel('Actual')
    plt.xlabel('Predicted')
    plt.show()
    print(f'Accuracy with {kernel_type} kernel: {accuracy:.4f}')
    print("\n")
evaluate_svm('rbf')
evaluate_svm('linear')

```



Accuracy with rbf kernel: 1.0000



Accuracy with linear kernel: 1.0000

Letter Recognition

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_score,
roc_curve
from sklearn.preprocessing import LabelBinarizer
import matplotlib.pyplot as plt
letters = pd.read_csv("letter-recognition.csv")
X = letters.iloc[:, 1:]
y = letters.iloc[:, 0] # assuming first column is label
```

```


y_binary = (y == 'A').astype(int) # modify based on actual dataset
X_train, X_test, y_train_bin, y_test_bin = train_test_split(X, y_binary,
test_size=0.2, random_state=42)
svm_model = SVC(kernel='rbf', probability=True)
svm_model.fit(X_train, y_train_bin)
y_pred = svm_model.predict(X_test)
y_prob = svm_model.predict_proba(X_test)[:, 1]
print("Letter Recognition Accuracy:", accuracy_score(y_test_bin, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test_bin, y_pred))
fpr, tpr, _ = roc_curve(y_test_bin, y_prob)
auc_score = roc_auc_score(y_test_bin, y_prob)
plt.figure()
plt.plot(fpr, tpr, label=f"ROC curve (AUC = {auc_score:.2f})")
plt.plot([0, 1], [0, 1], 'k--')
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend(loc="lower right")
plt.grid(True)
plt.show()

```

```

Letter Recognition Accuracy: 0.996
Confusion Matrix:
[[3850  1]
 [ 15 134]]

```



Horse Mule dataset

```

import pandas as pd
import matplotlib.pyplot as plt
from sklearn.svm import SVC
from sklearn.preprocessing import LabelEncoder

```

```


df = pd.read_csv("horse_mule_data.csv")
# Encode 'Horse'=0, 'Mule'=1
df['Label'] = LabelEncoder().fit_transform(df['Label'])
X = df[['Height', 'Weight']]
y = df['Label']
model = SVC(kernel='linear', C=1000) # High C -> fewer support vectors
model.fit(X, y)
support_vectors = model.support_vectors_
accuracy = model.score(X, y)
print("Accuracy:", accuracy)
print("Support Vectors:\n", support_vectors)
print("Number of Support Vectors:", len(support_vectors))
colors = ['red' if label == 0 else 'blue' for label in y]
plt.figure(figsize=(8,6))
plt.scatter(X['Height'], X['Weight'], c=colors, label='Data Points')
plt.scatter(support_vectors[:, 0], support_vectors[:, 1],
            s=200, facecolors='none', edgecolors='black', label='Support Vectors')
plt.xlabel("Height")
plt.ylabel("Weight")
plt.title("SVM Classification (3 Support Vectors)")
plt.legend()
plt.grid(True)
plt.show()

```

```

Accuracy: 1.0
Support Vectors:
[[145. 380.]
 [140. 350.]]
Number of Support Vectors: 2

```



Program 8

Implement Random forest ensemble method on a given dataset.

Screenshot:

15/04/25 LAB-8

[Random Forest]

Difference b/w Decision Tree and Random Forest:

	Decision Tree	Random Forest
1.	Single tree	ensemble of multiple trees
2.	High risk of overfitting	less prone to overfitting due to averaging multiple trees
3.	Can be less accurate	Generally more accurate
4.	Faster to train	Slower to train
5.	Poor handling of missing data	Better handling of missing data

Parameters of RandomForestClassifier():

- n_estimators - no of trees in forest
- criterion - function to measure quality of split
- max_depth - maximum depth of trees
- n_jobs - number of processors to use for parallel computation
- max_features - no of features to consider when splitting node
- Random_state - controls Randomness ensuring reproducibility

Algorithm:

1. Training dataset

• Randomly select sample replacement
• grow decision tree
• at each split, choose Random subset of features
• split using best features

3. Aggregate prediction
4. Output prediction

CODE:

```

import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

X = load_iris().data
y = load_iris().target

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

rf = RandomForestClassifier(n_estimators=10,
                           random_state=42)

rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy with n_estimators=10 is {accuracy}")

n_estimators_range = [10, 50, 100, 200, 500]
accuracy_scores = []

for n in n_estimators_range:
    rf = RandomForestClassifier(n_estimators=n,
                               random_state=42)

```

```

rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
accuracy_score.append(accuracy_score(y_test, y_pred))

best_n_estimators = n_estimators[accuracy_score.
                                index(max(accuracy_score))]

Output:
Accuracy with n_estimators=10 is 1.000
best n_estimators: 10 with accuracy: 1.000

Output: [train_size]
Training set size: 712 samples
Testing set size: 179 samples
Accuracy: 0.8212
Confusion Matrix:
[[ 91  14]
 [ 18  56]]

Output:
n_estimators = 30
Accuracy: 0.8101
AUC score: 0.8830
Confusion Matrix:
[[ 92  13]
 [ 21  53]]

```

```

n_estimators = 20
Accuracy: 0.8212
AUC score: 0.8857
Confusion Matrix:
[[ 91  14]
 [ 18  56]]

n_estimators = 30
Accuracy: 0.8101
AUC score: 0.8854
Confusion Matrix:
[[ 89  16]
 [ 18  56]]

```


Code:

Iris.csv

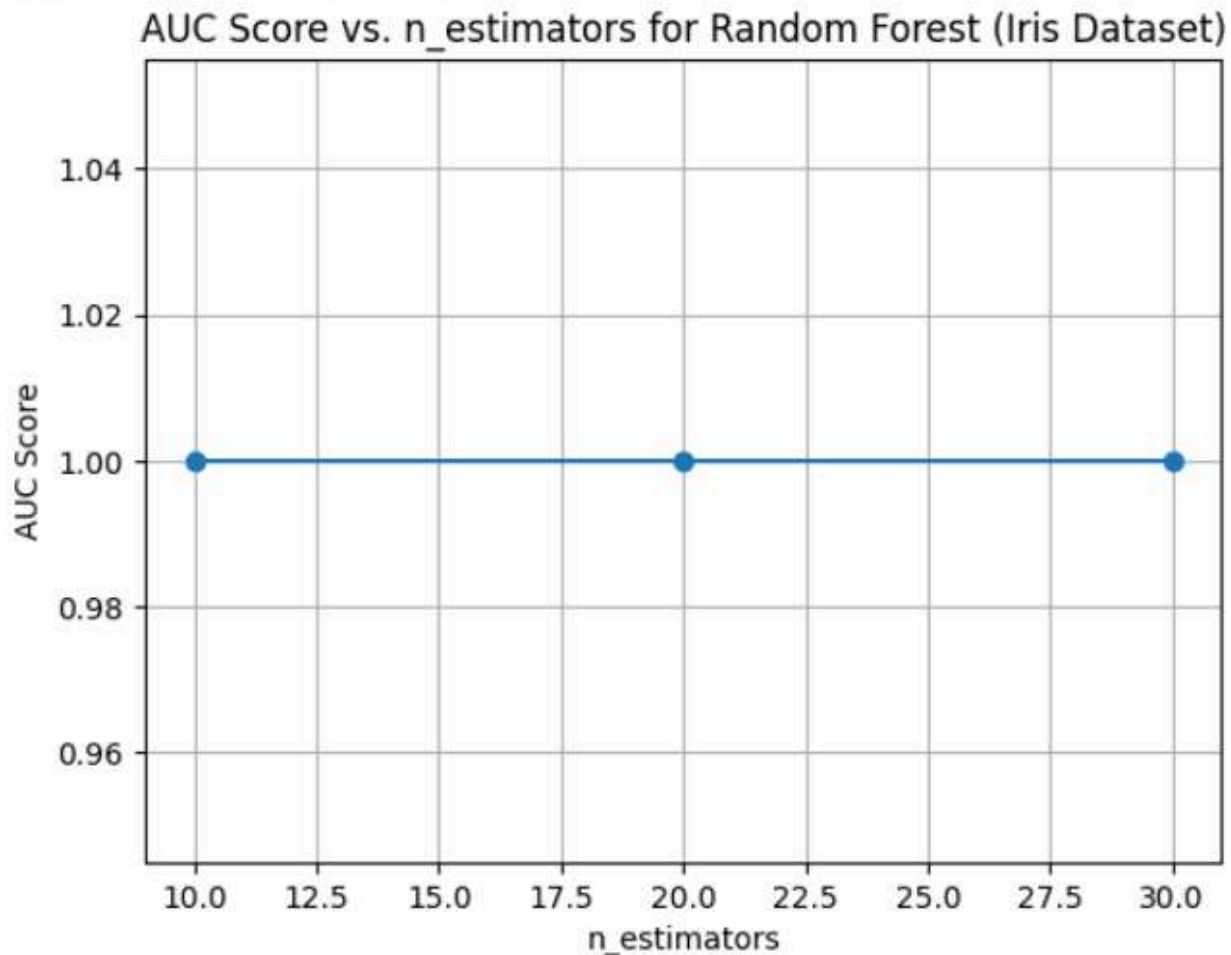
```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
iris = pd.read_csv("iris.csv")
X = iris.iloc[:, :-1]
y = iris.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
rf_classifier = RandomForestClassifier(random_state=42)
rf_classifier.fit(X_train, y_train)
y_pred = rf_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy with default n_estimators (10): {accuracy:.4f}")
best_accuracy = 0
best_n_estimators = 0
for n_estimators in range(10, 201, 10):
    rf_classifier = RandomForestClassifier(n_estimators=n_estimators, random_state=42)
    rf_classifier.fit(X_train, y_train)
    y_pred = rf_classifier.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_n_estimators = n_estimators
print(f"Best accuracy: {best_accuracy:.4f} achieved with n_estimators = {best_n_estimators}")
Accuracy with default n_estimators (10): 1.0000
Best accuracy: 1.0000 achieved with n_estimators = 10
```

```
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
from sklearn.preprocessing import label_binarize
from sklearn.multiclass import OneVsRestClassifier
iris = load_iris()
X, y = iris.data, iris.target
y = label_binarize(y, classes=[0, 1, 2])
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
n_estimators_values = [10, 20, 30]
auc_scores = []
for n_estimators in n_estimators_values:
    rf_classifier = OneVsRestClassifier(RandomForestClassifier(n_estimators=n_estimators,
random_state=42))
    rf_classifier.fit(X_train, y_train)
    y_pred_proba = rf_classifier.predict_proba(X_test)
```

```

auc_scores.append(roc_auc_score(y_test, y_pred_proba, average='weighted', multi_class='ovr'))
print(f'AUC Score for n_estimators = {n_estimators}: {auc_scores[-1]}')
plt.plot(n_estimators_values, auc_scores, marker='o')
plt.xlabel('n_estimators')
plt.ylabel('AUC Score')
plt.title('AUC Score vs. n_estimators for Random Forest (Iris Dataset)')
plt.grid(True)
plt.show()
AUC Score for n_estimators = 10: 1.0
AUC Score for n_estimators = 20: 1.0
AUC Score for n_estimators = 30: 1.0

```



Train.csv

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.preprocessing import LabelEncoder
df = pd.read_csv("train.csv")
df = df.drop(columns=['PassengerId', 'Name', 'Ticket', 'Cabin'])
df['Age'].fillna(df['Age'].median(), inplace=True)

```

```

df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
label_encoders = {}
for col in ['Sex', 'Embarked']:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
X = df.drop(columns=['Survived'])
y = df['Survived']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
print(f'Accuracy: {accuracy:.4f}')
print("Confusion Matrix:")
print(conf_matrix)

```

Accuracy: 0.8212

Confusion Matrix:

```

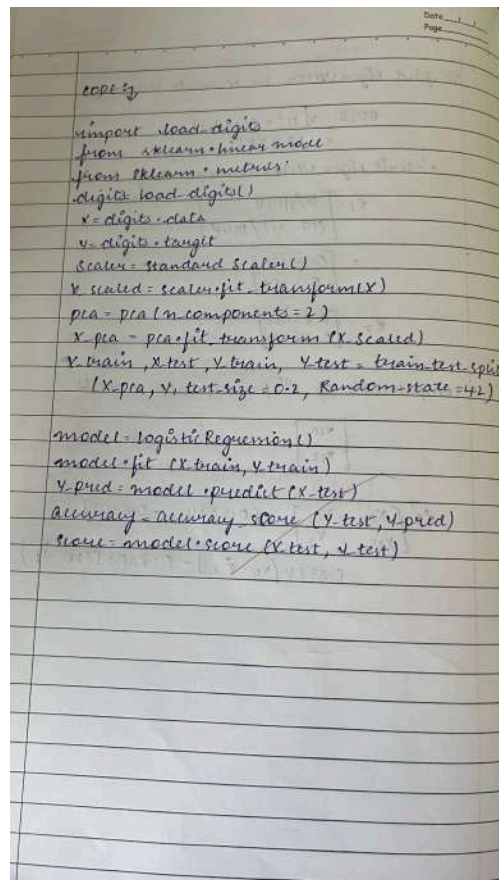
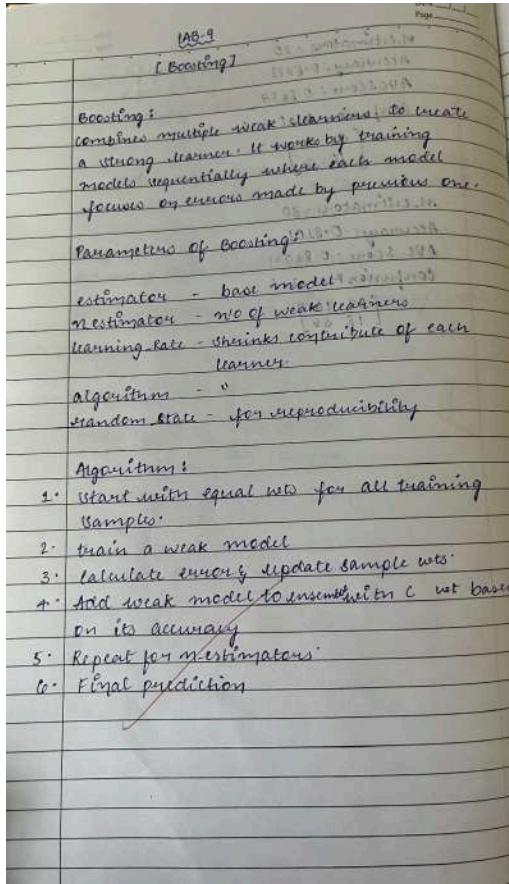
[[92 13]
 [19 55]]

```

Program 9

Implement Boosting ensemble method on a given dataset.

Screenshot:



Code:

Income.csv

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
data = pd.read_csv('income.csv')
X = data.drop('income_level', axis=1)
y = data['income_level']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
ada_boost = AdaBoostClassifier(n_estimators=50, random_state=42)
ada_boost.fit(X_train, y_train)
```

```


y_pred = ada_boost.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy}')
conf_matrix = confusion_matrix(y_test, y_pred)
print(f'Confusion Matrix:\n{conf_matrix}')
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=ada_boost.classes_, yticklabels=ada_boost.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - AdaBoost')
plt.tight_layout()
plt.show()

```

```

Accuracy: 0.8327362063670796
Confusion Matrix:
[[7003  411]
 [1223 1132]]

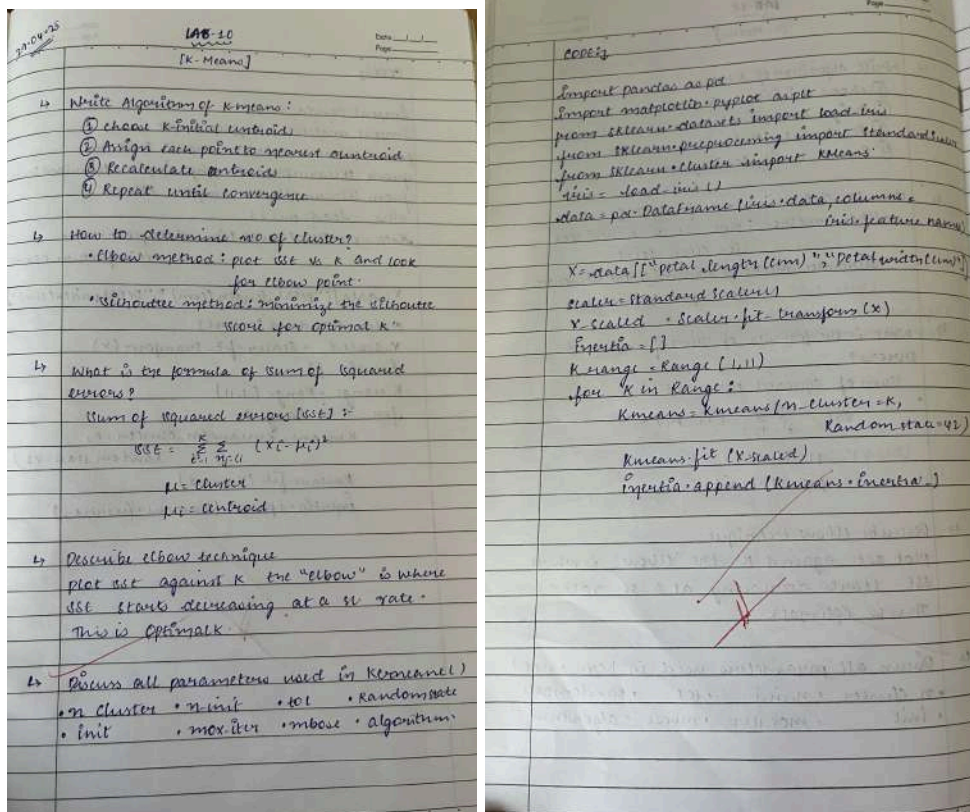
```



Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file.

Screenshot:



Code:


Income.csv

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import adjusted_rand_score
import matplotlib.pyplot as plt
import random
np.random.seed(42)
names = [f'Person_{i}' for i in range(1, 51)]
ages = np.random.randint(20, 60, 50)
incomes = np.random.randint(20000, 100000, 50)
df = pd.DataFrame({
    'Name': names,
    'Age': ages,
    'Income': incomes
})
df.to_csv('income.csv', index=False)
print("'income.csv' created successfully.")
```

```

from google.colab import files
files.download('income.csv')
data = pd.read_csv('income.csv')
X = data[['Age', 'Income']]
X_train, X_test = train_test_split(X, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
sse = []
k_range = range(1, 11)
for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_train_scaled)
    sse.append(kmeans.inertia_)
plt.plot(k_range, sse, marker='o')
plt.xlabel('Number of Clusters')
plt.ylabel('SSE (Inertia)')
plt.title('Elbow Method: SSE vs Number of Clusters')
plt.grid(True)
plt.show()
k = 3 # Change this based on the elbow plot
model = KMeans(n_clusters=k, random_state=42)
model.fit(X_train_scaled)
train_preds = model.predict(X_train_scaled)
test_preds = model.predict(X_test_scaled)
true_labels_train = [random.randint(0, k-1) for _ in range(len(train_preds))]
accuracy = adjusted_rand_score(true_labels_train, train_preds)
print("Adjusted Rand Index (proxy accuracy):", round(accuracy, 2))

```





Iris.csv

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
iris = load_iris()
data = pd.DataFrame(iris.data, columns=iris.feature_names)
X = data[["petal length (cm)", "petal width (cm)"]]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
inertia = []
K_range = range(1, 11)
for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_scaled)
    inertia.append(kmeans.inertia_)
plt.figure(figsize=(8, 5))
plt.plot(K_range, inertia, marker='o')
plt.title("Elbow Method for Optimal k")
plt.xlabel("Number of clusters (k)")
```

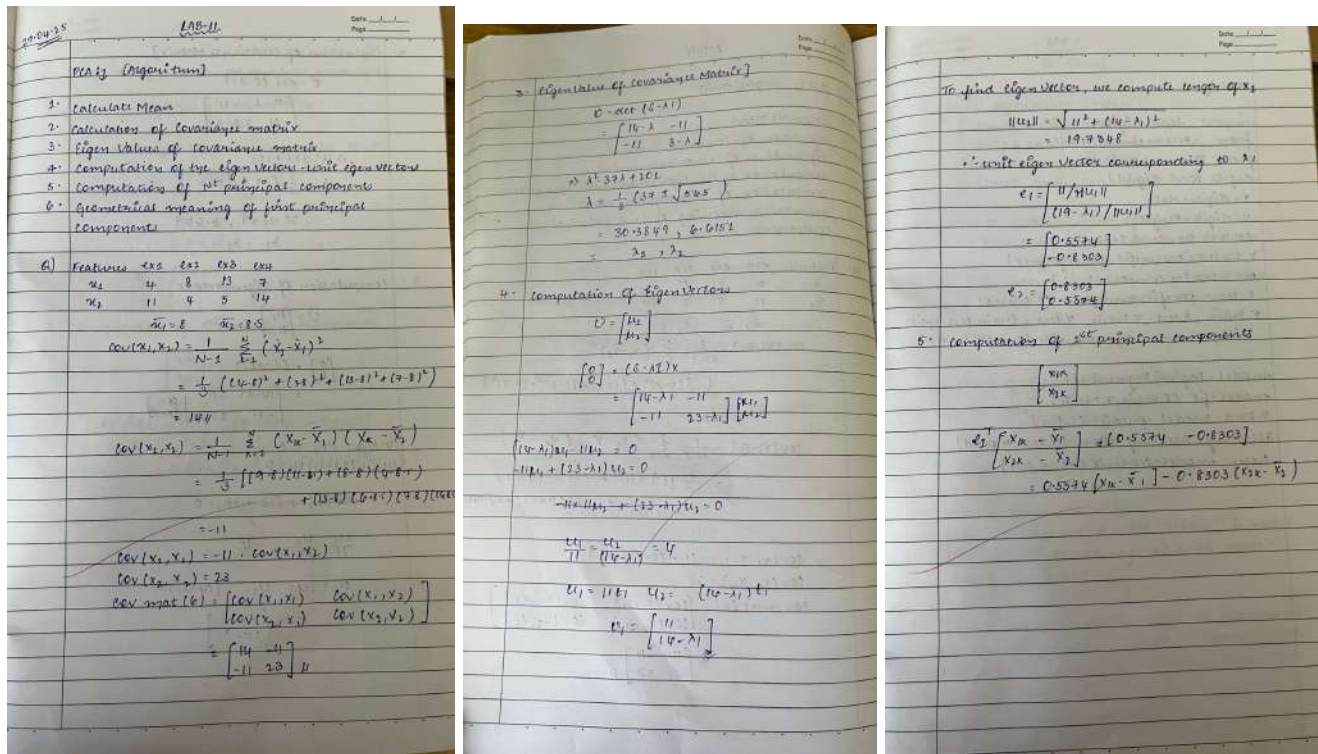


```
plt.ylabel("Inertia")  
plt.grid(True)  
plt.show()
```



Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method.
Screenshot:



Code:

```
from sklearn.datasets import load_digits
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

digits = load_digits()
X = digits.data
y = digits.target

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

X_train, X_test, y_train, y_test = train_test_split(
    X_pca, y, test_size=0.2, random_state=42
)

model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
score = model.score(X_test, y_test)
print("Model Score (accuracy using .score()):", round(score, 4))
```

```
print(" Accuracy using PCA with 2 components:", round(accuracy, 4))
```

```
Model Score (accuracy using .score()): 0.5389
```

```
Accuracy using PCA with 2 components: 0.5389
```

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA
from scipy.stats import zscore
df = pd.read_csv("/content/heart.csv") # Adjust path if needed
z_scores = np.abs(zscore(df.select_dtypes(include=[np.number])))
df = df[(z_scores < 3).all(axis=1)]
df_encoded = df.copy()
for col in df_encoded.select_dtypes(include=["object"]).columns:
    if df_encoded[col].nunique() <= 2:
        le = LabelEncoder()
        df_encoded[col] = le.fit_transform(df_encoded[col])
    else:
        df_encoded = pd.get_dummies(df_encoded, columns=[col], drop_first=True)
X = df_encoded.drop("target", axis=1) # Replace 'target' if it's named differently
y = df_encoded["target"]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC()
}
print("Model Accuracies (without PCA):")
for name, model in models.items():
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    acc = accuracy_score(y_test, preds)
    print(f"{name}: {acc:.4f}")
pca = PCA(n_components=0.95) # Retain 95% variance
X_pca = pca.fit_transform(X_scaled)
X_train_pca, X_test_pca, _, _ = train_test_split(X_pca, y, test_size=0.2, random_state=42)
print("\nModel Accuracies (with PCA):")
for name, model in models.items():
    model.fit(X_train_pca, y_train)
    preds = model.predict(X_test_pca)
```

```
acc = accuracy_score(y_test, preds)
print(f'{name}: {acc:.4f}')
```

Model Accuracies (without PCA):

Logistic Regression: 0.8103

Random Forest: 0.7759

SVM: 0.7931

Model Accuracies (with PCA):

Logistic Regression: 0.8276

Random Forest: 0.8103

SVM: 0.7586